Lecture 5

ARMA Models

Lars A. Lochstoer

UCLA Anderson School of Management

Winter 2020

Outline

Autoregressive Models

ARMA Models

- **parsimonious** description of (univariate) time series (mimicking autocorrelation etc.)
- very useful tools for forecasting (and commonly used in industry)
  - forecasting sales, earnings revenue growth at the firm level or at the industry level
  - forecasting GPD growth, inflation at the national level

Autoregressive process of order 1

- lagged returns might be useful in predicting returns.
- we consider a model that allows for this:

$$r_{t+1} = \phi_0 + \phi_1 r_t + \varepsilon_{t+1}, \qquad \varepsilon_{t+1} \sim \mathrm{WN}(0, \sigma_\varepsilon^2)$$

  - $\{\varepsilon_t\}$ represents the 'news':

$$\varepsilon_t = r_t - E_{t-1}[r_t]$$

    ε_t is what you know about the process at t but not at $t-1$
  - Economists often call ε_t the 'shocks' or 'innovations'.
- this model is referred to as an **AR(1)**

Transition density

Definition

Given an information set $\mathcal{F}_t$, the **transition density** of a random variable r_{t+1} is the conditional distribution of r_{t+1} given by:

$$r_{t+1} \sim p(r_{t+1}|\mathcal{F}_t; \boldsymbol{\theta})$$

- The information set $\mathcal{F}_t$ is often (but not always) the history of the process $r_t, r_{t-1}, r_{t-2}, \ldots$
- In this case, the transition density is written:

$$r_{t+1} \sim p(r_{t+1}|r_t, r_{t-1}, \ldots, ; \boldsymbol{\theta})$$

- A transition density is **Markov** if it depends on its finite past.

AR(1) transition density

- Consider the AR(1) model with Gaussian shocks

$$r_{t+1} = \phi_0 + \phi_1 r_t + \varepsilon_{t+1}, \qquad \varepsilon_t \sim \mathrm{N}(0, \sigma_\varepsilon^2)$$

- The transition density is **Markov of order 1**.

$$r_{t+1} \sim p(r_{t+1} | r_t; \boldsymbol{\theta})$$

the rest of the history $r_{t-2}, r_{t-3}, \ldots$ is irrelevant.

- With Gaussian shocks ε_t, the transition density is:

$$r_{t+1} \ \sim \ \mathrm{N}(\phi_0 + \phi_1 r_t, \sigma_\varepsilon^2)$$

- conditional mean and conditional variance:

$$\begin{aligned} E\left[r_{t+1} | r_t\right] &= \phi_0 + \phi_1 r_t, \\ V\left[r_{t+1} | r_t\right] &= V\left[\varepsilon_{t+1}\right] = \sigma_\varepsilon^2. \end{aligned}$$

Unconditional mean of AR(1)

- assume that the series is covariance-stationary
- compute the unconditional mean μ.
  - take unconditional expectations:

$$E\left[r_{t+1}\right] = \phi_0 + \phi_1 E\left[r_t\right].$$

  - use stationarity: $E\left[r_{t+1}\right] = E\left[r_t\right] = \mu$:

$$\mu = \phi_0 + \phi_1 \mu,$$

    and solving for the unconditional mean:

$$\mu = \frac{\phi_0}{1 - \phi_1}.$$

- mean exists if $\phi_1 \neq 1$ and is zero if $\phi_0 = 0$

Mean Reversion

- if $\phi_1 \neq 1$, we can rewrite the AR(1) process as:

$$r_{t+1} - \mu = \phi_1 (r_t - \mu) + \varepsilon_{t+1}.$$

- suppose $0 < \phi_1 < 1$
  - when $r_t > \mu$, the process is expected to get **closer** to the mean:

$$E_t[r_{t+1} - \mu] = \phi_1 (r_t - \mu) < (r_t - \mu).$$

  - when $r_t < \mu$, the process is expected to get **closer** to the mean:

$$E_t[r_{t+1} - \mu] = \phi_1 (r_t - \mu) > (r_t - \mu).$$

- the smaller ϕ_1, the higher the speed of mean reversion

Mean Reversion

- we can rewrite the AR(1) process as:

$$r_{t+2} - \mu = \phi_1^2 \left(r_t - \mu\right) + \phi_1 \varepsilon_{t+1} + \varepsilon_{t+2}.$$

- suppose $0 < \phi_1 < 1$
  - when $r_t > \mu$, the process is expected to get **closer** to the mean:

$$E_t[r_{t+2} - \mu] = \phi_1^2 \left(r_t - \mu\right) < \left(r_t - \mu\right).$$

  - when $r_t < \mu$, the process is expected to get **closer** to the mean:

$$E_t[r_{t+2} - \mu] = \phi_1^2 \left(r_t - \mu\right) > \left(r_t - \mu\right).$$

Half Life

- we can rewrite the AR(1) process as:

$$r_{t+h} - \mu = \phi^h (r_t - \mu) + \phi^{h-1} \varepsilon_{t+1} + \ldots + \varepsilon_{t+h}.$$

- suppose $0 < \phi_1 < 1$
  - at the **half-life**, the process is expected to cover **1/2** of the distance to the mean:

$$E_t[r_{t+h} - \mu] = \phi_1^h (r_t - \mu) = .5 (r_t - \mu).$$

- the half-life is defined by setting $\phi_1^h = 0.5$ and solving

$$h = \log(0.5) / \log(\phi_1)$$

Variance of AR(1)

Compute the unconditional variance:

- take the expectation of the square of:

$$r_{t+1} - \mu = \phi_1 \left(r_t - \mu \right) + \varepsilon_{t+1}.$$

- we obtain the following expression for the unconditional variance:

$$V \left[r_{t+1} \right] = \frac{\sigma_\varepsilon^2}{1 - \phi_1^2},$$

  provided that $\phi_1^2 < 1$ because the variance has to be positive and bounded

- covariance stationarity requires that

$$-1 < \phi_1 < 1.$$

- in addition, if $-1 < \phi_1 < 1$, we can show that the series is covariance stationary because the mean and variance are finite

Continuous-Time Model

Definition

In a continuous-time model, the log of stock prices, $p_t = \log P_t$, follows an **Ornstein-Uhlenbeck process** if:

$$dp_t \;=\; \kappa(\mu_p - p_t)dt + \sigma_p dB_t \tag{1}$$

Continuous-time version of a discrete-time, Gaussian AR(1) process.

Suppose we observe the process (1) at discrete intervals Δt, then this is equivalent to:

$$p_t \;=\; \mu + \phi_1(p_{t-1} - \mu) + \sigma\varepsilon_t \qquad \varepsilon_t \sim \mathrm{N}(0,1)$$

where

- $\phi_1 \;=\; \exp(-\kappa\Delta t)$
- $\mu \;=\; \mu_p$
- $\sigma^2 \;=\; (1 - \exp(-2\kappa\Delta t))\frac{\sigma_p^2}{2\kappa}$.

Dynamic Multipliers

- use the expression for the mean of the AR(1) to obtain:

$$r_{t+1} - \mu = \phi_1 \left(r_t - \mu \right) + \varepsilon_{t+1}.$$

- by repeated substitution, we get:

$$r_t - \mu = \sum_{i=0}^{t} \phi_1^i \varepsilon_{t-i} + \phi^{t+1} \left(r_{-1} - \mu \right).$$

- value of r_t at t is stated as a function of the **history of shocks** $\left\{ \varepsilon_\tau \right\}_{\tau=0}^{\tau=t}$ and its value at time $t = -1$

- effect of shocks die out over time provided that $-1 < \phi_1 < 1$.

Dynamic Multipliers

calculate the effect of a change ε_0 on r_t:

$$\frac{\partial [r_t - \mu]}{\partial \varepsilon_0} = \phi_1^t.$$

$$\frac{\partial [r_{t+j} - \mu]}{\partial \varepsilon_t} = \phi_1^j.$$

in a covariance stationary model, *dynamic multiplier* only depends on j, not on t

Again, note that we need $|\phi_1| < 1$ for a stationary (non-explosive) system where shocks die out: $\lim_{j \to \infty} \phi_1^j = 0$

MA(infinity) representation

- use the expression for the mean of the AR(1) to obtain:

$$r_{t+1} - \mu = \phi_1 (r_t - \mu) + \varepsilon_{t+1}.$$

- by repeated substitution:

$$r_t - \mu = \sum_{i=0}^{\infty} \phi_1^i \varepsilon_{t-i}.$$

  - **linear** function of past innovations!
  - fits into class of linear time series

Autocovariances of an AR(1)

- take the unconditional expectation:

$$\left(r_t-\mu\right)\left(r_{t-j}-\mu\right)=\phi_1\left(r_{t-1}-\mu\right)\left(r_{t-j}-\mu\right)+\varepsilon_t\left(r_{t-j}-\mu\right).$$

- this yields:

$$E\left[\left(r_t-\mu\right)\left(r_{t-j}-\mu\right)\right]=\phi_1 E\left[\left(r_{t-1}-\mu\right)\left(r_{t-j}-\mu\right)\right]+E\left[\varepsilon_t\left(r_{t-j}-\mu\right)\right].$$

- or, using notation from Lecture 9:

$$\begin{aligned}\gamma_j &= \phi_1\gamma_{j-1}, \qquad j>0\\ \gamma_0 &= \phi_1\gamma_{-1}+\sigma_\varepsilon^2, \qquad j=0\end{aligned}$$

- note that $\gamma_{-j}=\gamma_j$

Autocorrelation Function

- it immediately implies that the ACF is:

$$\rho_j = \phi_1 \rho_{j-1}, \qquad j \geq 0$$

  and $\rho_0 = 1$

- combing these two equations imply that:

$$\rho_j = \phi_1^j$$

  - exponential decay at a rate ϕ_1

Autocorrelation Function of an AR(1)

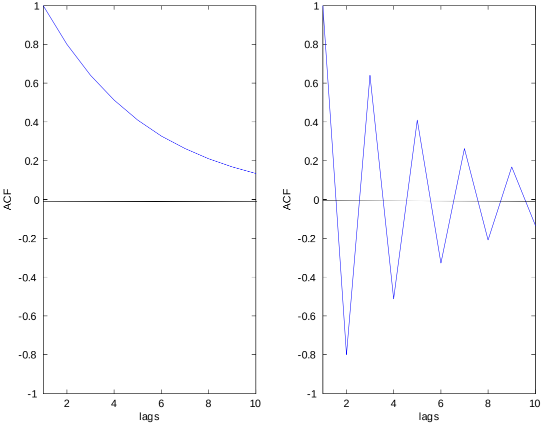

Autocorrelation Function for AR(1). The left panel considers $\phi_1 = 0.8$. The right panel considers $\phi_1 = -0.8$.

AR(p)

Definition

The $\mathbf{AR}(p)$ model is defined as:
$r_t = \phi_0 + \phi_1 r_{t-1} + \ldots + \phi_p r_{t-p} + \varepsilon_t, \qquad \varepsilon_t \sim \text{WN}\left(0, \sigma_\varepsilon^2\right)$

- other lagged returns might be useful in predicting returns
- similar to multiple regression model with p lagged variables as explanatory variables
- the $\mathbf{AR}(p)$ is **Markov of order p**.

Conditional Moments

- conditional mean and conditional variance:

$$
\begin{aligned}
E\left[r_{t+1}|r_t, ..., r_{t-p+1}\right] &= \phi_0 + \phi_1 r_t + \ldots + \phi_p r_{t-p+1} \\
V\left[r_{t+1}|r_t, ..., r_{t-p+1}\right] &= V\left[\varepsilon_{t+1}\right] = \sigma_\varepsilon^2
\end{aligned}
$$

- moments conditional on $r_t, \ldots, r_{t-p+1}$ are not correlated with $r_{t-i}, i \geq p$

AR(2)

- consider the model:

$$r_t = \phi_0 + \phi_1 r_{t-1} + \phi_2 r_{t-2} + \varepsilon_t \qquad \varepsilon_t \sim \text{WN}\left(0, \sigma_\varepsilon^2\right)$$

- take unconditional expectations to compute the mean

$$E\left[r_t\right] = \phi_0 + \phi_1 E\left[r_{t-1}\right] + \phi_2 E\left[r_{t-2}\right]$$

- Assuming stationarity and solving for the mean:

$$E\left[r_t\right] = \mu = \frac{\phi_0}{1 - \phi_1 - \phi_2}$$

provided that $\phi_1 + \phi_2 \neq 1$.

- using this expression for μ write the model in deviation from means:

$$r_t - \mu = \phi_1\left(r_{t-1} - \mu\right) + \phi_2\left(r_{t-2} - \mu\right) + \varepsilon_t$$

Autocorrelations of an AR(2)

- take the expectation of :

$$\begin{aligned}(r_t - \mu)\left(r_{t-j} - \mu\right) &= \phi_1 (r_{t-1} - \mu)\left(r_{t-j} - \mu\right) \\ &\quad + \phi_2 (r_{t-2} - \mu)\left(r_{t-j} - \mu\right) + \varepsilon_t \left(r_{t-j} - \mu\right)\end{aligned}$$

- this yields:

$$\begin{aligned}E\left[(r_t - \mu)\left(r_{t-j} - \mu\right)\right] &= \phi_1 E[(r_{t-1} - \mu)\left(r_{t-j} - \mu\right)] \\ &\quad + \phi_2 E[(r_{t-2} - \mu)\left(r_{t-j} - \mu\right)] \\ &\quad + E\left[\varepsilon_t \left(r_{t-j} - \mu\right)\right]\end{aligned}$$

- or, using different notation:

$$\begin{aligned}\gamma_j &= \phi_1 \gamma_{j-1} + \phi_2 \gamma_{j-2}, & j > 0 \\ \gamma_0 &= \phi_1 \gamma_{-1} + \phi_2 \gamma_{-2} + \sigma_\varepsilon^2, & j = 0\end{aligned}$$

Autocorrelations of an AR(2)

- the ACF:

$$\begin{aligned}\rho_j &= \phi_1\rho_{j-1} + \phi_2\rho_{j-2}, \qquad j \geq 2 \\ \rho_0 &= \phi_1\rho_{-1} + \phi_1\rho_{-2} + \sigma_\varepsilon^2/\gamma_0, \qquad j = 0\end{aligned}$$

which implies that the ACF of an AR(2) satisfies a second-order difference equation:

$$\begin{aligned}\rho_1 &= \phi_1\rho_0 + \phi_2\rho_1 \\ \rho_j &= \phi_1\rho_{j-1} + \phi_2\rho_{j-2}, \qquad j \geq 2\end{aligned}$$

Roots

Definition

The second-order difference equation for the ACF:

$$(1 - \phi_1 B - \phi_2 B^2)\rho_j = 0,$$

where B is the **back-shift operator**: $B\rho_j = \rho_{j-1}$

Note that we can write the above as:

$$(1 - \varpi_1 B)(1 - \varpi_2 B)\rho_j = 0$$

- A useful factorization
- Intuitively, the AR(2) is an "AR(1) on top of another AR(1)"
- From AR(1) math, we had that each AR(1) is stationary if its autocorrelation is less than one in absolute value.
- The 'roots' ϖ_j should satisfy similar property for AR(2) to be stationary

Finding the roots

A simple case:

$$
\begin{aligned}
1 - \phi_1 B - \phi_2 B^2 &= (1 - \varpi_1 B)(1 - \varpi_2 B) \\
&= 1 - (\varpi_1 + \varpi_2) B + \varpi_1 \varpi_2 B^2
\end{aligned}
$$

and so we solve using the relations:

$$
\begin{aligned}
\phi_1 &= \varpi_1 + \varpi_2 \\
\phi_2 &= -\varpi_1 \varpi_2
\end{aligned}
$$

The solutions to this are the inverses to the solutions to the second order polynomial in the scalar-valued x:

$$(1 - \phi_1 x - \phi_2 x^2) = 0,$$

- the solutions to this equation are given by:

$$x_1, x_2 = \frac{\phi_1 \pm \sqrt{\phi_1^2 + 4\phi_2}}{-2\phi_2}$$

- the inverses are the **characteristic roots**: $\varpi_1 = x_1^{-1}$ and $\varpi_2 = x_2^{-1}$

Roots (real, distinct case)

- two characteristic roots: $\varpi_1 = x_1^{-1}$ and $\varpi_2 = x_2^{-1}$
- both characteristic roots are real-valued if the discriminant is greater than zero: $\phi_1^2 + 4\phi_2 > 0$
  - then we can factor the polynomial as:

$$(1 - \phi_1 B - \phi_2 B^2) = (1 - \varpi_1 B)(1 - \varpi_2 B)$$

  - *two AR(1) models on top of each other*
- The ACF will decay like an AR(1).

Roots (complex-valued case)

- two characteristic roots: $\varpi_1 = x_1^{-1}$ and $\varpi_2 = x_2^{-1}$
- both characteristic roots are complex-valued if the discriminant is negative: $\phi_1^2 + 4\phi_2 < 0$
- Then, $\varpi_1 = x_1^{-1}$ and $\varpi_2 = x_2^{-1}$ are complex numbers.
- The ACF will look like damped sine and cosine waves.

Autocorrelation for AR(2)

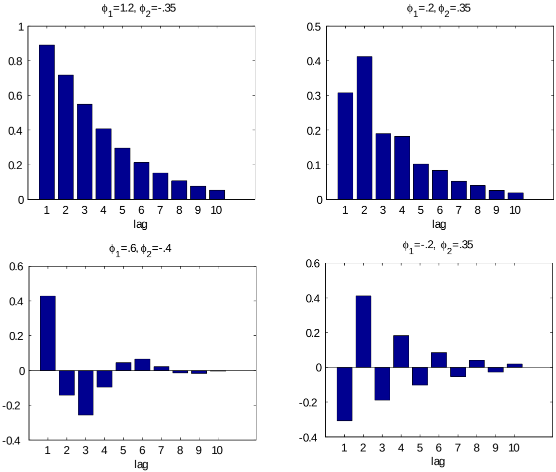

Autocorrelation Function for AR(2) processes.

AR(2) Example: The Dividend Price Ratio

- The stock market Dividend to Price ratio is:
  - Sum of last year's dividends to firms in the market divided by current market value
  - A "Valuation Ratio"
  - Very slow-moving (persistent); quarterly postWW2 data for U.S.:

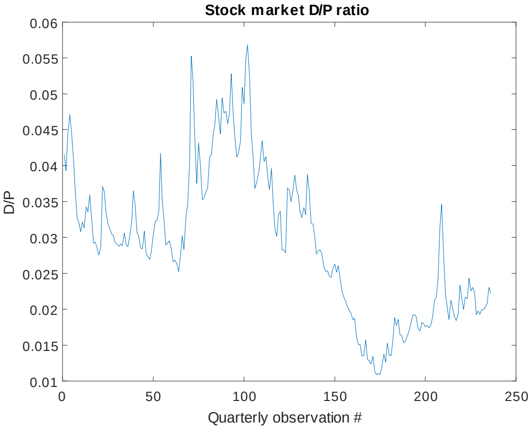

Estimate AR(2) on this variable

ARIMA(2,0,0) Model:

Conditional Probability Distribution: Gaussian

Parameter	Value	Standard Error	t Statistic
Constant	0.00123254	0.00074679	1.65045
AR{1}	1.09319	0.0527929	20.7072
AR{2}	-0.137308	0.051282	-2.67752
Variance	7.84588e-06	1.92026e-07	40.8583

- Stationarity test:

$$1 - 1.09319x + 0.13731x^2 = 0$$

  - Roots greater than 1, so stationary despite $\phi_1 = 1.093 > 1$ as $\phi_2 = -0.137$.
  - Unconditional mean:

$$\mu = \frac{0.00123254}{1 - 1.09319 + 0.13731} = 0.0279$$

AR(2) DP prediction

Pred_DP1 = uncond_mean + phi1*(DP(2:end)-uncond_mean) + phi2*(DP(1:end-1)-uncond_mean);

Pred_DP2 = uncond_mean + phi1*(Pred_DP1-uncond_mean) + phi2*(DP(2:end)-uncond_mean);

Pred_DP3 = uncond_mean + phi1*(Pred_DP2-uncond_mean) + phi2*(Pred_DP1-uncond_mean);

etc.

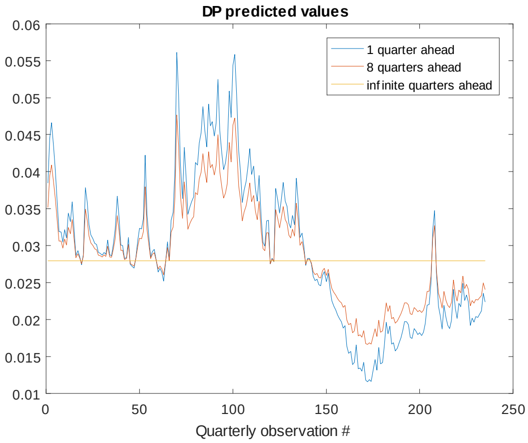

Stationarity

- Recall: The modulus of $z = a + bi$ is $|z| = \sqrt{a^2 + b^2}$. Thus, for real numbers the modulus is simply the absolute value.

Result:

- An AR(1) process is stationary if its characteristic root is less than one, i.e. if $1/x = \phi_1$ is less than one in modulus. This condition implies that $\rho_j = \phi_1^j$ converges to zero as $j \to \infty$.

- An AR(2) process is stationary if the two characteristic roots ϖ_1 and ϖ_2 (the inverses of the solutions to those two equations) are less than one in modulus.

Stationarity of AR(p)

- **An AR(p) process is stationary if all p characteristic roots of the below polymonial are less than one in modulus**

$$1 - \phi_1 x - \phi_2 x^2 - ... - \phi_p x^p = 0$$

- see chapter 2 in Hamilton (1994) for details.

Partial Autocorrelation Function

Definition

The PACF of a stationary series is defined as $\{\phi_{j,j}\}, j = 1, \ldots, n$

$$
\begin{aligned}
r_t &= \phi_{0,1} + \boldsymbol{\phi}_{1,1} r_{t-1} + v_{1t} \\
r_t &= \phi_{0,2} + \phi_{1,2} r_{t-1} + \boldsymbol{\phi}_{2,2} r_{t-2} + v_{2t} \\
r_t &= \phi_{0,3} + \phi_{1,3} r_{t-1} + \phi_{2,3} r_{t-2} + \boldsymbol{\phi}_{3,3} r_{t-3} + v_{3t} \\
& \ldots
\end{aligned}
$$

- These are simple multiple regressions that can be estimated with least squares.
- $\phi_{p,p}$ shows the incremental contribution of r_{t-p} to r_t over an $AR(p-1)$ model

PACF

Definition

The **sample partial autocorrelations (PACF)** of a time series are defined as $\widehat{\phi}_{1,1}, \widehat{\phi}_{2,2}, \ldots, \widehat{\phi}_{p,p}, \ldots,$

Partial Autocorrelation Function

The PACF of an $AR(p)$ satisfies:

1. $\hat{\phi}_{p,p} \longrightarrow \phi_p$ as sample size increases
2. $\hat{\phi}_{j,j} \longrightarrow 0$ for $j > p$

- for an AR(p) series, the sample PACF cuts off after lag p
- $\Rightarrow$ look at the sample PACF to determine an appropriate value of p

PACF of Daily Log Returns

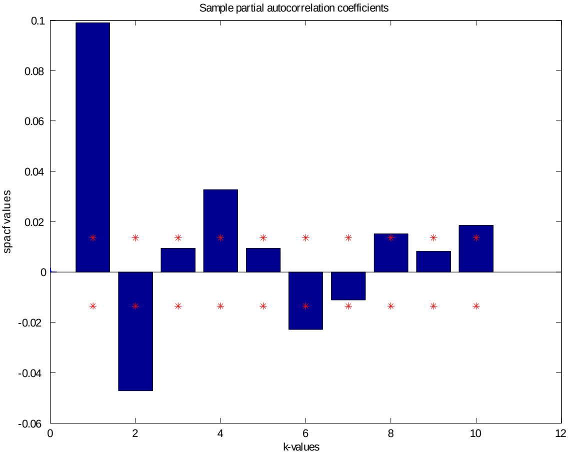

PACF for Daily log Returns on VW-CRSP Index. Two standard error bands around zero. 1926-2007.

Information Criteria

- information criteria help determine the **optimal lag length**
- the Akaike (1973) information criterion:

$$AIC = -2\ln(likelihood) + 2(number\ of\ parameters)$$

- the Bayesian information criterion of Schwarz (1978):

$$BIC = -2\ln(likelihood) + \ln T(number\ of\ parameters)$$

  - the BIC penalty depends on the sample size T
- for different values of p, compute $AIC(p)$ and/or $BIC(p)$ pick the lag length with the minimum AIC/BIC

Manufacturing White Noise

- to check the performance of the AR model you've selected: **check the residuals**!!
- residuals should look like **white noise**
    - look at the ACF of the residuals
    - perform Ljung-Box test on residuals
    - $Q(m) \sim \chi^2(m-p)$ where p is the lag length of the $AR(p)$ model

Forecasting

- suppose we have an AR(p) model
- we want to forecast r_{t+h} using all the info $\mathcal{F}_t$ available at t
- assume we choose the forecast to minimize the **mean square error**:

$$E\left[\left(y - y_{prediction}\right)^2\right]$$

- The conditional mean minimizes the mean squared forecast error.
- we will come back to **optimal forecasting** later

1-step ahead forecast error

- the $AR(p)$ model is given by:

$$r_{t+1} = \phi_0 + \phi_1 r_t + \ldots + \phi_p r_{t-p+1} + \varepsilon_{t+1}$$

- take the conditional expectation:

$$E_t\left[r_{t+1}\right] = \phi_0 + \phi_1 r_t + \ldots + \phi_p r_{t-p+1}$$

- the one-step ahead forecast error:

$$v_t(1) = r_{t+1} - \phi_0 - \sum_{i=1}^{p} \phi_i r_{t-i+1} = \varepsilon_{t+1}$$

- the variance of the one-step ahead forecast error:

$$V\left[v_t\left(1\right)\right] = \sigma_\varepsilon^2$$

  - if ε_t is normally distributed, then the 95 % confidence interval:

$$\pm 1.96\sigma_\varepsilon$$

2-step ahead forecast error

- the $AR(p)$ model is given by:

$$r_{t+2} = \phi_0 + \phi_1 r_{t+1} + \ldots + \phi_p r_{t-p+2} + \varepsilon_{t+2}$$

- we just take the conditional expectation:

$$E_t\left[r_{t+2}\right] = \phi_0 + \phi_1 \widehat{r}_t(1) + \ldots + \phi_p r_{t-p+2}$$

- the two-step ahead forecast error:

$$v_t(2) = \phi_1 v_t(1) + \varepsilon_{t+2} = \phi_1 \varepsilon_{t+1} + \varepsilon_{t+2}$$

- the variance of the two-step ahead forecast error:

$$\boxed{V\left[v_t\left(2\right)\right] = \sigma_\varepsilon^2 (1 + \phi_1^2)}$$

  - the variance of the two-step ahead forecast error is larger than the variance of the one-step ahead forecast error

Multi-step ahead forecast error

Result:

The h-step ahead forecast is given by:

$$\hat{r}_t\left(h\right) = \phi_0 + \sum_{i=1}^{p} \phi_i \hat{r}_t\left(h-i\right)$$

where $\hat{r}_t\left(j\right) = r_{t+j}$ if $j < 0$.

- the h-step ahead forecast converges to the unconditional expectation $E(r_t)$ as $h \longrightarrow \infty$
- this is referred to as **mean reversion**

Estimation: conditional least squares

- assume we observe or can condition on the first p observations.
- AR(p) model is then a linear regression model:

$$r_t = \phi_0 + \phi_1 r_{t-1} + \ldots + \phi_p r_{t-p} + \varepsilon_t, \qquad t = p+1, \ldots, T$$

- using least squares, the fitted model is

$$\widehat{r}_t = \widehat{\phi}_0 + \widehat{\phi}_1 r_{t-1} + \ldots + \widehat{\phi}_p r_{t-p}$$

and the residual is $v_t = r_t - \widehat{r}_t$

- the estimated variance of the residuals is:

$$\widehat{\sigma}^2_\varepsilon = \frac{\sum_{t=p+1}^{T} v_t^2}{T - 2p - 1}$$

ML Estimation

- alternatively, we could use maximum likelihood.
- the log-likelihood function is:

$$\ln p(r_1, r_2, \ldots, r_T; \boldsymbol{\theta}) = \sum_{t=2}^{T} \ln p(r_t | r_{t-1}, \ldots, r_1; \boldsymbol{\theta}) + \ln p(r_1; \boldsymbol{\theta})$$

- for example, assume Gaussian shocks ε_t then $p(r_t | r_{t-1}, \ldots, r_{t-p}; \boldsymbol{\theta})$ is normal
- the difference between least squares and ML estimation of $(\phi_0, \phi_1, \ldots, \phi_p)$ are the initial distributions $p(r_1; \boldsymbol{\theta}), p(r_2 | r_1; \boldsymbol{\theta}) \ldots$.
- Conditional least squares of an AR(p) drops the first p terms in the likelihood.

Example: ML Estimation of AR(1)

- assume the initial value r_1 comes from the stationary dist.
- unconditional moments:

$$E[r_1] = \frac{\phi_0}{1-\phi_1}, \qquad V[r_1] = \frac{\sigma_\varepsilon^2}{1-\phi_1^2},$$

- hence, the density $p(r_1;\theta)$ of the first observation r_1 is normal with the above (unconditional) mean and variance
- for $t > 1$, the conditional moments:

$$E[r_t|r_{t-1}] \;=\; \phi_0 + \phi_1 r_{t-1}, \qquad V[r_t|r_{t-1}] \;=\; \sigma_\varepsilon^2$$

- hence, the conditional density $p(r_t|r_{t-1};\boldsymbol{\theta})$ is normal with the above (conditional) mean and (conditional) variance

Example: ML Estimation of AR(1)

- the log-likelihood function is:

$$\begin{aligned}
\ln p(r_1, r_2, \ldots, r_T; \boldsymbol{\theta}) &= \sum_{t=2}^{T} \ln p(r_t | r_{t-1}, \ldots, r_1; \boldsymbol{\theta}) + \ln p(r_1; \boldsymbol{\theta}) \\
&= -\frac{1}{2} \sum_{t=2}^{T} \left(\ln(2\pi) + \ln(\sigma_\varepsilon^2) + \frac{(r_t - \phi_0 - \phi_1 r_{t-1})^2}{\sigma_\varepsilon^2} \right) \\
&\quad + \ln p(r_1; \boldsymbol{\theta})
\end{aligned}$$

- choose parameters $\boldsymbol{\theta} = (\phi_0, \phi_1, \sigma_\varepsilon^2)$ to maximize the log-likelihood function

$p(r_1; \boldsymbol{\theta})$ is typically chosen to be the stationary distribution

$$r_1 \sim \mathrm{N}\left(\frac{\phi_0}{1 - \phi_1}, \frac{\sigma_\varepsilon^2}{1 - \phi_1^2}\right)$$

Exact vs. Conditional ML

- the conditional ML estimator drops the initial condition
- exact log-likelihood function:

$$\ln p(r_1, r_2, \ldots, r_T; \boldsymbol{\theta}) = \sum_{t=2}^{T} \ln p(r_t | r_{t-1}, \ldots, r_1; \boldsymbol{\theta}) + \ln p(r_1 | \boldsymbol{\theta})$$

- conditional log-likelihood function:

$$\ln p(r_{p+1}, \ldots, r_T; \boldsymbol{\theta}) = \sum_{t=p+1}^{T} \ln p(r_t | r_{t-1}, \ldots, r_1; \boldsymbol{\theta})$$

- the conditional log-likelihood 'conditions' on the first data point and drops the first p terms.
- Conditional ML is the same as least squares. The solution can be calculated analytically.

Summary: AR(p) models

- *dynamic model*, e.g. AR(p):

$$r_t = \phi_0 + \phi_1 r_{t-1} + \ldots + \phi_p r_{t-p} + \varepsilon_t$$

  - constant determines mean through: $\mu = \frac{\phi_0}{1-\phi_1-\phi_2-\ldots-\phi_p}$
  - coefficients $(\phi_1, \phi_2, \ldots, \phi_p)$ must satisfy stationarity restrictions for a well-specified model:
  - objective: parsimonious model of dynamics of r_t

- For AR(p) models, you can maximize the conditional MLE in closed-form...conditional least squares....but there is no guarantee that it will satisfy the stationarity restrictions.

- Calculating the full MLE requires numerical optimization.

Application:

Bond Pricing

Bond Notation

- an n-period zero coupon bond pays one dollar n periods from now
- notation:
  - $P_t^{(n)}$ denotes the price of an n-period zero-coupon bond.
  - $p_t^{(n)} = \log(P_t^{(n)})$ denotes the log price
  - the yield of an n-period zero-coupon bond is:

$$y_t^{(n)} \equiv -\frac{1}{n} p_t^{(n)}$$

  - the holding period return is:

$$hpr_{t+1}^{(n)} \equiv p_{t+1}^{(n-1)} - p_t^{(n)}$$

  - the short term interest rate $y_t^{(1)}$ is given special notation r_t

Term Structure of Interest Rates

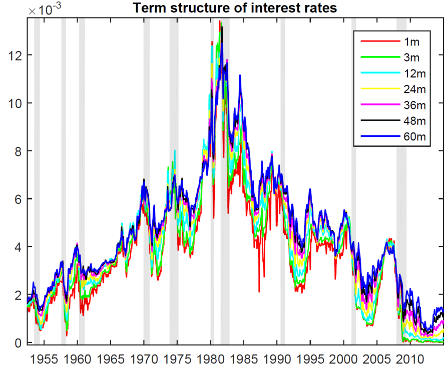

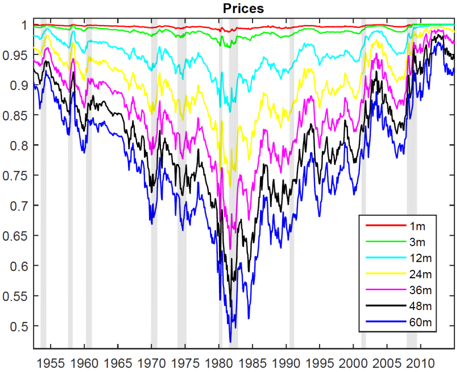

CRSP Fama-Bliss Zero-Coupon Bond Data. Sample: 1952.6-2014.12. Monthly data. Left: yields. Right: prices

ACF and PACF of 1 month yield

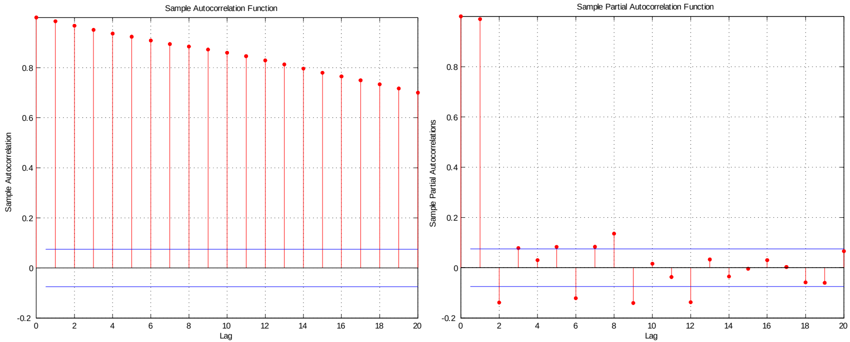

CRSP Fama-Bliss Zero-Coupon Bond Data. Sample: 1952.6-2014.12. Monthly data. Yield on one month zero-coupon bond $y_t^{(1)}$.

- ACF is persistent. PACF drops off after 1 month.

Bond Pricing: Expectations Hypothesis

- discrete time models of bond pricing
- examine the simplest possible model: a single factor model, investors risk-neutral with respect to interest rate risks
- the single factor g_t follows an AR(1):

$$g_{t+1} = (1-\phi)\mu + \phi g_t + \sigma \varepsilon_{t+1}, \quad \varepsilon_{t+1} \sim \mathrm{N}(0,1)$$

- In "exponential-affine" term structure models the log risk-free rate is linear in factors

$$r_t \; = \; \delta_0 + \delta_1' g_t$$

- Since only one factor in our example, might as well let log short rate be the factor $r_t = y_t^{(1)} = g_t$ (ie, $\delta_0 = 0$ and $\delta_1 = 1$)

Pricing of zero-coupon default-free bonds

- A zero-coupon bond only pays off principal (normalize this to \$1), no coupons.
- If investors are risk-neutral, they always discount with the risk-free rate.
- so, the price of a 1-period zero-coupon bond would be

$$P_t^{(1)} = e^{-r_t} \times \$1$$

- the price of a 2-period zero-coupon bond would be:

$$P_t^{(2)} = E_t[e^{-r_t} P_{t+1}^{(1)}] = E_t[e^{-r_t} e^{-r_{t+1}}]$$

- Price of *n*-period bond

$$P_t^{(n)} = E_t[\exp\left(-\sum_{j=0}^{n-1} r_{t+j}\right)]$$

Pricing (cont'd)

- We need to take expectations of AR(1) variable in exponential
- Use the fact:
$$E\left[e^{x}\right] = e^{\mu+\frac{1}{2}\sigma^2}$$
if $x \sim N\left(\mu, \sigma^2\right)$
- Consider 2-period bond

$$\begin{aligned}
P_t^{(2)} &= e^{-r_t} E_t\left[e^{-r_{t+1}}\right] \\
&= e^{-r_t} E_t\left[e^{-((1-\phi)\mu+\phi r_t+\sigma\varepsilon_{t+1})}\right] \\
&= e^{-r_t} e^{-(1-\phi)\mu-\phi r_t} E_t\left[e^{-\sigma\varepsilon_{t+1}}\right] \\
&= e^{-2\mu-(1+\phi)(r_t-\mu)} E_t\left[e^{-\sigma\varepsilon_{t+1}}\right] \\
&= e^{-2\mu-(1+\phi)(r_t-\mu)+\frac{1}{2}\sigma^2}.
\end{aligned}$$

Price of an n-period bond

The Price of an *n*-period zero-coupon bond is:

$$P_t^{(n)} = E_t\left[e^{-r_t} P_{t+1}^{(n-1)}\right]$$

Solving recursively, we get:

$$P_t^{(n)} = \exp\left(a_n + b_n r_t\right)$$

where

$$\begin{aligned} a_n &= a_{n-1} + b_{n-1}\left(1-\phi\right)\mu + \frac{1}{2}\sigma^2 b_{n-1}^2, \\ b_n &= b_{n-1}\phi - 1, \end{aligned}$$

with initial conditions $a_1 = 0$ and $b_1 = -1$.

- This implies yields are:

$$y_t^{(n)} = \tilde{a}_n + \tilde{b}_n g_t \qquad \tilde{a}_n = -\frac{1}{n} a_n \qquad \tilde{b}_n = -\frac{1}{n} b_n$$

- a_n and b_n are difference equations fit through the *cross-section* of yields.

Estimation

- monthly U.S. zero coupon Fama-Bliss data from CRSP.
- 1,3,12,24,36,48,60 month yields
- Model estimated by maximum likelihood (add noise to yield observations)
- short rate $r_t = g_t$ is the one-month yield $y_t^{(1)}$

Parameter	Model	Sample moment	Model-implied
μ	$mean(y_t^{(1)})$	0.00355	0.00362
$\frac{\sigma^2}{1-\phi^2}$	$var(y_t^{(1)})$	5.976e-06	8.368e-06
ϕ	$\rho_{12}(y_t^{(1)})$	0.9756	0.9895

- All yields perfectly conditionally correlated due to 1-factor structure
- High current interest rates expected to revert to mean so long-term rates lower than short-term rates.
- Average slope of yield curve is slightly below zero (since investors' risk-neutral and there is a convexity term), contrary to data

Multiple Factors

- the Vasicek (1977) adds "risk prices" to the this bare-bones model, but is still a simple, one-factor model
- this model cannot capture the slope or curvature of yields, only the level of interest rates.
- we need a richer model with more factors, where we let g_t be a vector of factors $\Rightarrow$ vector autoregressive process
  - PCA of yields indicated we need three factors to explain yields.
- Fixed Income class will discuss more on this

Moving Average Models

AR(infinity)

- in theory the true data generating process could be an AR(∞):

$$x_t = \phi_0 + \phi_1 x_{t-1} + \phi_2 x_{t-2} + \ldots + \varepsilon_t$$

- implementation:
  - **infinite** number of parameters
- solution: constrain parameters

$$x_t = \phi_0 - \theta_1 x_{t-1} - \theta_1^2 x_{t-2} - \theta_1^3 x_{t-3} - \ldots + \varepsilon_t$$

where $\phi_i = -\theta_1^i, i \geq 1$

AR(infinity) to MA(1)

- solution: constrain parameters

$$x_t + \theta_1 x_{t-1} + \theta_1^2 x_{t-2} + \theta_1^3 x_{t-3} + \ldots = \phi_0 + \varepsilon_t$$

- this can be restated as an MA(1) model:

$$x_t = \phi_0(1 - \theta_1) + (1 - \theta_1 B)\varepsilon_t$$

  - MA(1) is a 'cheap' version of an AR(∞).

- general form of MA(1) model is:

$$x_t = \mu + (1 - \theta_1 B)\varepsilon_t = \mu + \varepsilon_t - \theta_1 \varepsilon_{t-1}$$

MA(q)

Definition

A **moving average process of order** q or MA(q) model is:

$$x_t = \mu + (1 - \theta_1 B - \ldots - \theta_q B^q)\varepsilon_t,$$

where $q > 0$

Stationarity

- consider the MA(1) model:

$$x_t = \mu + (1 - \theta_1 B)\varepsilon_t.$$

- compute the variance of an MA(1) model:

$$V\left[\mu + (1 - \theta_1 B)\,\varepsilon_t\right] = (1 + \theta_1^2)\sigma_\varepsilon^2.$$

- compute the variance of an MA(q) model:

$$V\left[\mu + (1 - \theta_1 B - ... - \theta_q B^q)\,\varepsilon_t\right] = (1 + \theta_1 + \theta_2^2 + \ldots + \theta_q^q)\sigma_\varepsilon^2.$$

Computing Autocovariances for MA(1)

- assume the unconditional mean $\mu = 0$
- pre-multiply the MA(1) model by r_{t-l}:

$$r_{t-j} r_t = r_{t-j} \varepsilon_t - \theta_1 r_{t-j} \varepsilon_{t-1}$$

- take expectations
- compute the auto-covariance of an MA(1) model:

$$\gamma_1 = -\theta_1 \sigma_\varepsilon^2, \ \gamma_j = 0, \ j > 1$$

- this implies the autocorrelations are:

$$\rho_1 = \frac{-\theta_1}{1 + \theta_1^2}, \ \rho_j = 0, \ j > 1$$

  - **the ACF is cut off after 1 lag**!

Computing Autocovariances for MA(2)

- the same argument implies that the autocorrelations of an MA(2) are:

$$\rho_1 = \frac{-\theta_1 + \theta_1\theta_2}{1 + \theta_1^2 + \theta_2^2}, \qquad \rho_2 = \frac{-\theta_2}{1 + \theta_1^2 + \theta_2^2} \qquad \rho_j = 0,\ j > 2$$

  - the ACF is cut off after 2

Forecasting with MA(1)

- consider an MA(1) model:

$$r_{t+1} = \mu + \varepsilon_{t+1} - \theta_1 \varepsilon_t$$

- take conditional expectations:

$$\widehat{r}_t(1) = E_t\left[r_{t+1}\right] = \mu - \theta_1 \varepsilon_t$$

$$\widehat{r}_t(2) = E_t\left[r_{t+2}\right] = \mu$$

- the one-step ahead forecast error is given by:

$$v_t(1) = r_{t+1} - \widehat{r}_t(1) = \varepsilon_{t+1}$$

  - the variance of the one-step ahead forecast error is σ_ε^2

Forecasting with MA(1)

- consider an MA(1) model:

$$r_{t+1} = \mu + \varepsilon_{t+1} - \theta_1 \varepsilon_t$$

- take conditional expectations:

$$\widehat{r}_t(1) = E_t\left[r_{t+1}\right] = \mu - \theta_1 \varepsilon_t$$

$$\widehat{r}_t(2) = E_t\left[r_{t+2}\right] = \mu$$

- the two-step ahead forecast error is given by:

$$v_t(2) = r_{t+2} - \widehat{r}_t(2) = \varepsilon_{t+2} - \theta_1 \varepsilon_{t+1}$$

  - the variance of the two-step ahead forecast error is $(1+\theta_1^2)\sigma_\varepsilon^2$
  - this is the unconditional variance

Forecasting with MA(2)

- consider an MA(2) model:

$$r_{t+1} = \mu + \varepsilon_{t+1} - \theta_1 \varepsilon_t - \theta_2 \varepsilon_{t-1}$$

- take conditional expectations:

$$\widehat{r}_t(1) = E_t[r_{t+1}] = \mu - \theta_1 \varepsilon_t - \theta_2 \varepsilon_{t-1}$$

$$\widehat{r}_t(2) = E_t[r_{t+2}] = \mu - \theta_2 \varepsilon_t$$

$$\widehat{r}_t(3) = E_t[r_{t+3}] = \mu$$

- the one-step ahead forecast error is given by:

$$v_t(1) = r_{t+1} - \widehat{r}_t(1) = \varepsilon_{t+1}$$

  - the variance of the one-step ahead forecast error is σ_ε^2

Forecasting with MA(2)

- consider an MA(2) model:

$$r_{t+1} = \mu + \varepsilon_{t+1} - \theta_1 \varepsilon_t - \theta_2 \varepsilon_{t-1}$$

- take conditional expectations:

$$\widehat{r}_t(1) = E_t\left[r_{t+1}\right] = \mu - \theta_1 \varepsilon_t - \theta_2 \varepsilon_{t-1}$$

$$\widehat{r}_t(2) = E_t\left[r_{t+2}\right] = \mu - \theta_2 \varepsilon_t$$

$$\widehat{r}_t(3) = E_t\left[r_{t+3}\right] = \mu$$

- the two-step ahead forecast error is given by:

$$v_t(2) = r_{t+2} - \widehat{r}_t(2) = \varepsilon_{t+2} - \theta_1 \varepsilon_{t+1}$$

  - the variance of the two-step ahead forecast error is $(1 + \theta_1^2)\sigma_\varepsilon^2$
  - this is smaller than the unconditional variance

Forecasting with MA(2)

- consider an MA(2) model:

$$r_{t+1} = \mu + \varepsilon_{t+1} - \theta_1 \varepsilon_t - \theta_2 \varepsilon_{t-1}$$

- take conditional expectations:

$$\widehat{r}_t(1) = E_t\left[r_{t+1}\right] = \mu - \theta_1 \varepsilon_t - \theta_2 \varepsilon_{t-1}$$

$$\widehat{r}_t(2) = E_t\left[r_{t+2}\right] = \mu - \theta_2 \varepsilon_t$$

$$\widehat{r}_t(3) = E_t\left[r_{t+3}\right] = \mu$$

- the three-step ahead forecast error is given by:

$$v_t(3) = r_{t+3} - \widehat{r}_t(3) = \varepsilon_{t+3} - \theta_1 \varepsilon_{t+2} - \theta_2 \varepsilon_{t+1}$$

  - the variance of the three-step ahead forecast error is $(1 + \theta_1^2 + \theta_2^2)\sigma_\varepsilon^2$
  - this is the unconditional variance

Maximum Likelihood

- MA(q) models can't be estimated using (conditional) least squares because the parameters are a non-linear function of the data

- MA(q) models are commonly estimated using Maximum Likelihood
- this involves assuming a parametric distribution for the shocks ε_t.
- Often, we assume ε_t are normally distributed.

ML Estimation of MA(1)

- conditional moments:

$$\begin{aligned} V\left[r_t | r_{t-1}\right] &= \sigma_\varepsilon^2, \\ E\left[r_t | r_{t-1}\right] &= \mu - \theta\varepsilon_{t-1} \end{aligned}$$

- hence, the density $p(r_t|\mathcal{F}_{t-1};\boldsymbol{\theta})$ of the first observation is normal with the above (conditional) mean and variance

- suppose we assume that $\varepsilon_0 = 0$.
- then $\varepsilon_1 = r_1 - \mu$
- then $\varepsilon_2 = r_2 - \mu - \theta_1\varepsilon_1$
- we can recursively calculate the sequence $\{\varepsilon_1, \varepsilon_2, \varepsilon_3, \ldots, \varepsilon_t\}$

ML Estimation of MA(1)

- hence, the log-likelihood function is:

$$
\begin{aligned}
\ln p(r_1, r_2, \ldots, r_T; \boldsymbol{\theta}) &= \sum_{t=2}^{T} \ln p(r_t | r_{t-1}, \ldots, r_1; \boldsymbol{\theta}) + \ln p(r_1; \boldsymbol{\theta}) \\
&= -\frac{1}{2} \sum_{t=1}^{T} \left(\ln(2\pi) + \ln(\sigma_\varepsilon^2) + \frac{(-\varepsilon_t)^2}{\sigma_\varepsilon^2} \right) \\
&\quad + \ln p(r_1; \boldsymbol{\theta})
\end{aligned}
$$

- choose parameters $\boldsymbol{\theta} = (\mu, \theta_1, \sigma_\varepsilon^2)$ to maximize the log-likelihood function

ACF and PACF

- *ACF* is useful for determining MA lag length:
  - autocorrelations are cut off at q for an $MA(q)$: $ACF(k) = 0$ for $k > q$
- *PACF* is useful for determining AR lag length
  - partial autocorrelations are cut off at p for an $AR(p)$: $PACF(k) = 0$ for $k > p$

ARMA Models

ARMA(p,q)

- certain processes can only be described by AR or MA models if we include lots of lags
  - unappealing (need to estimate lots of parameters)
- natural solution: ARMA(p, q) processes

ARMA(p,q)

- consider an ARMA$(1,1)$ model:

$$r_t - \phi_1 r_{t-1} = \phi_0 + \varepsilon_t - \theta_1 \varepsilon_{t-1} \qquad \varepsilon_t \sim \mathrm{WN}(0, \sigma_\varepsilon^2)$$

  with $\theta_1 \neq \phi_1$

- the unconditional mean of an ARMA$(1,1)$ has the same expression as an AR(1)

$$E\left[r_t\right] = \frac{\phi_0}{1-\phi_1}$$

- we can re-write the process as:

$$r_t - \mu = \phi_1 (r_{t-1} - \mu) + \varepsilon_t - \theta_1 \varepsilon_{t-1}$$

- take expectations of $[r_t - \mu]^2$ to compute the variance:

$$V\left[r_t\right] = \phi_1^2 V\left[r_{t-1}\right] + \sigma_\varepsilon^2 + \theta_1^2 \sigma_\varepsilon^2 - 2\phi_1\theta_1 E\left[\varepsilon_{t-1}\left(r_{t-1} - \mu\right)\right]$$

ARMA(1,1)

- this reduces to:

$$V[r_t] = \phi_1^2 V[r_t] + \sigma_\varepsilon^2 + \theta_1^2\sigma_\varepsilon^2 - 2\phi_1\theta_1\sigma_\varepsilon^2$$

- collecting terms, we get:

$$V[r_t] = \sigma_\varepsilon^2 \frac{1+\theta_1^2 - 2\phi_1\theta_1}{1-\phi_1^2}$$

- obviously, we need $\phi_1^2 < 1$
  - same stationarity requirement as for AR(1)

ACF of ARMA(1,1)

- to compute the auto-covariances:

$$
\begin{aligned}
E\left[(r_t-\mu)\left(r_{t-j}-\mu\right)\right] &= \phi_1 E\left[(r_{t-1}-\mu)\left(r_{t-j}-\mu\right)\right] \\
&\quad + E\left[\varepsilon_t\left(r_{t-j}-\mu\right)\right] \\
&\quad - \theta_1 E\left[\varepsilon_{t-1}\left(r_{t-j}-\mu\right)\right]
\end{aligned}
$$

- for $j=1$, we get: $\gamma_1 = \phi_1\gamma_0 - \theta_1\sigma_\varepsilon^2$
- this implies that the ACF is given by:

$$
\begin{aligned}
\rho_1 &= \phi_1 - \theta_1\frac{\sigma_\varepsilon^2}{\gamma_0} \\
\rho_j &= \phi_1\rho_{j-1}, \qquad j>1
\end{aligned}
$$

Autocorrelation for ARMA(1,1)

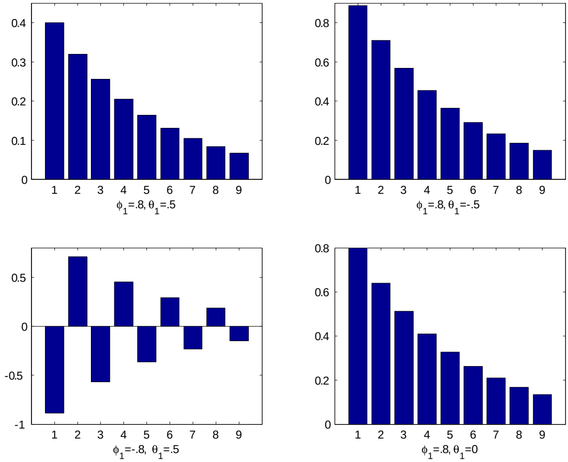

Autocorrelation Function for ARMA(1,1) processes.

PACF of ARMA(1,1)

- PACF does not die out at some lag
- slow decay (as is the case for MA models)

ARMA(p,q)

- consider an ARMA(p, q) model:

$$r_t = \phi_0 + \sum_{i=1}^{p} \phi_i r_{t-i} + \varepsilon_t - \sum_{i=1}^{q} \theta_i \varepsilon_{t-i}, \qquad \varepsilon_t \sim \mathrm{WN}(0, \sigma_\varepsilon^2)$$

- using the backshift operator

$$(1 - \phi_1 B - \ldots - \phi_p B^p) r_t = \phi_0 + (1 - \theta_1 B - \ldots - \theta_q B^q) \varepsilon_t$$

D/P autocorrelation function

- Revisiting the D/P ratio
  - Sample autocorrelation function:

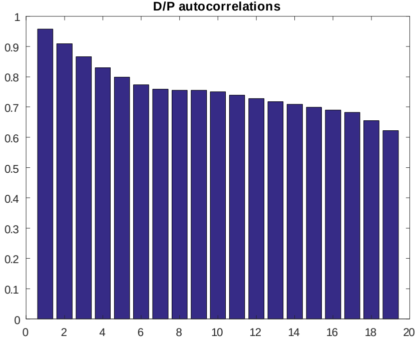

- 'Drop off' for about first 4 lags, then stable...
  - Indicates a 4 lags of MA might be a good representation + 1 lag AR

D/P as ARMA(1,4)

```
ARIMA(1,0,4) Model:
   --------------------
   Conditional Probability Distribution: Gaussian
```

Parameter	Value	Standard Error	t Statistic
Constant	0.000484308	0.000550722	0.879406
AR{1}	0.980784	0.0158269	61.9696
MA{1}	0.103568	0.0573102	1.80715
MA{2}	-0.172191	0.0611522	-2.81577
MA{3}	-0.148333	0.0632631	-2.34469
MA{4}	-0.106098	0.0582711	-1.82076
Variance	7.50796e-06	1.64215e-07	45.7203

- Forecast by getting sample series of residuals, then plug in as needed for forecasts

$$\mu = \frac{0.000484}{1 - 0.9807}$$

$$\begin{aligned}
E_t\left[DP_{t+1}\right] &= \mu + 0.98\left(DP_t\text{-}\mu\right) + 0.10\varepsilon_t\text{-}0.17\varepsilon_{t-1}\text{-}0.14\varepsilon_{t-2}\text{-}0.11\varepsilon_{t-3}, \\
E_t\left[DP_{t+2}\right] &= E_t\left[E_{t+1}\left[DP_{t+2}\right]\right] \\
&= E_t\left[\mu + 0.98\left(DP_{t+1}\text{-}\mu\right) + 0.10\varepsilon_{t+1}\text{-}0.17\varepsilon_t\text{-}0.14\varepsilon_{t-1}\text{-}0.11\varepsilon_{t-2}\right] \\
&= \mu + 0.98\left(E_t\left[DP_{t+1}\right]\text{-}\mu\right)\text{-}0.17\varepsilon_t\text{-}0.14\varepsilon_{t-1}\text{-}0.11\varepsilon_{t-2}
\end{aligned}$$

etc.

MA representation

- start from this expression:

$$(1-\phi_1 B - \ldots - \phi_p B^p) r_t = \phi_0 + (1 - \theta_1 B - \ldots - \theta_q B^q)\varepsilon_t$$

- re-arranging this expression delivers an *MA* representation:

$$r_t = \frac{\phi_0}{(1-\phi_1 B - \ldots - \phi_p B^p)} + \frac{(1-\theta_1 B - \ldots - \theta_q B^q)}{(1-\phi_1 B - \ldots - \phi_p B^p)}\varepsilon_t$$

- more succinctly:

$$r_t = \mu + \psi(B)\varepsilon_t$$

- stationarity: the solutions of $(1 - \phi_1 x - \ldots - \phi_p x^p) = 0$ should lie outside of the unit circle

Impulse-Response Function

- consider the MA representation:

$$r_t = \mu + \psi(B)\varepsilon_t$$

- this can be written out as:

$$r_t = \mu + \varepsilon_t + \psi_1\varepsilon_{t-1} + \psi_2\varepsilon_{t-2} + \dots$$

  where $\{\psi_i\}$ is the *impulse response* function of the ARMA model.
- the coefficients $\{\psi_i\}$ are functions of the parameters $\{\phi_i\}$ and $\{\theta_i\}$
- the impulse response function shows the effect today of a shock k periods ago:

$$\frac{\partial r_t}{\partial \varepsilon_{t-k}} = \psi_k$$

Forecasting

- consider the MA representation:

$$r_t = \mu + \psi(B)\varepsilon_t$$

- the h-period ahead forecast :

$$\widehat{r}_t(h) = \mu + \psi_h \varepsilon_t + \psi_{h+1}\varepsilon_{t-1} + \dots$$

- the h-period ahead forecast error can be stated as:

$$v_t(h) = \varepsilon_{t+h} + \psi_1 \varepsilon_{t+h-1} + \psi_2 \varepsilon_{t+h-2} + \dots + \psi_{h-1}\varepsilon_{t+1}$$

- the variance of the h-step ahead forecast error is:

$$V\left[v_t\left(h\right)\right] = \left(1 + \psi_1^2 + \psi_2^2 + \dots + \psi_{h-1}^2\right)\sigma_\varepsilon^2$$

Variance of Forecast Error

- the variance of the *h*-step ahead forecast error is:

$$V\left[v_t\left(h\right)\right] = \left(1 + \psi_1^2 + \psi_2^2 + \ldots + \psi_2^h\right)\sigma_\varepsilon^2$$

  - non-decreasing function of forecast horizon
  - variance of forecast error converges to variance of process

$$V\left[v_t\left(h\right)\right] \longrightarrow V\left[r_t\right]$$

  as $h \longrightarrow \infty$

References

Akaike, H. (1973). Information theory and an extension of the maximum likelihood principle. In B. Petrov and F. Cs aki (Eds.), 2nd International Symposium on Information Theory, Tsahkadsor, Armenia, USSR, pp. 267-281. Budapest: Akademiai Kiado.

Ang, A. and M. Piazzesi (2003). A no-arbitrage vector autoregression of term structure dynamics with macroeconomic and latent variables. Journal of Monetary Economics 50, 745-787.

Box, G. E. P. and G. Jenkins (1970). Time Series Analysis: Forecasting and Control. San Francisco, CA: Holden-Day.

Hamilton, J. D. (1994). Time Series Analysis. Princeton, NJ: Princeton University Press.

Ljung, G. M. and G. E. P. Box (1978). On a measure of a lack of t in time series models. Biometrika 65(2), 297-303.

Schwarz, G. E. (1978). Estimating the dimension of a model. The Annals of Statistics 6(2), 461-464.

Vasicek, O. A. (1977). An equilibrium characterization of the term structure. Journal of Financial Economics 5(2), 177-188.